

ENGR 4350: Applied Deep Learning

Logistic Regression: Part 1

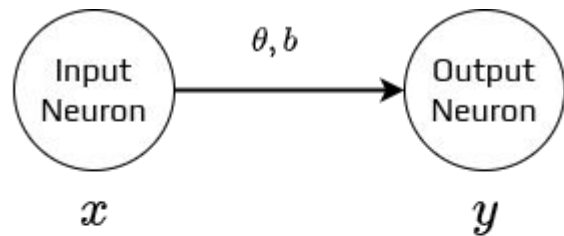
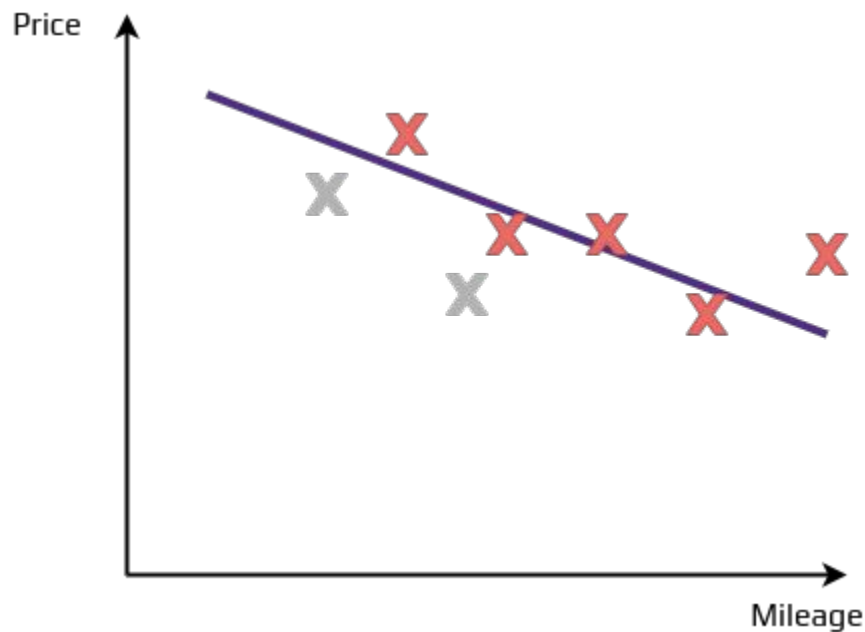
08/31/2022



Outline

- An example of neural network
- Logistic Regression
 - Forward Pass
 - Loss Function
 - Gradient Descent


A Neural Network Example



$$y = \sigma(\theta x + b)$$



A Neural Network Example

Make	Model	Year	Mileage	MPG	Buying Price	Selling Price
Ford	Edge	2018	50,000	23	\$19,000	\$10,000
Toyota	Land Cruiser	2020	10,000	15	\$80,000	\$50,000
VW	Golf	2010	150,000	36	\$7,000	\$2,000



$$\vec{y} = \sigma(\vec{\theta}\vec{x} + \vec{b})$$

A Neural Network Example



Binary Classification

- Complex decision makings can be simplified to classification problems.
 - Vehicle control
 - Robotic arm control
 - Gaming
 - ...
- Binary classification works most of the time.
 - Visitor identification
 - Animal protection
 - Farming
 - ...

Logistic Regression

Logistic regression estimates the probability of an event occurring, $P(\mathbf{y} = \mathbf{1} | \mathbf{x})$ such as voted or didn't vote, based on a given dataset of independent variables. Since the outcome is a probability, the dependent variable is bounded between 0 and 1.

- Classification
- Prediction
- Rating

...

E.g. Given the “make, model, year, mileage, MPG” of vehicles, estimate probabilities of prices of these vehicles under \$20,000.

Problem Settings

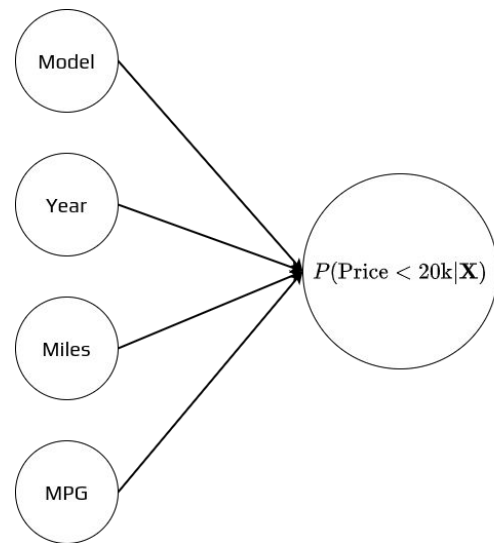
Dataset: $\left\{ \left(\mathbf{x}^{(1)}, y^{(1)} \right), \left(\mathbf{x}^{(2)}, y^{(2)} \right), \dots, \left(\mathbf{x}^{(m)}, y^{(m)} \right) \right\}$

$$\text{Features: } \mathbf{X} = \begin{bmatrix} x_1^{(1)} & x_2^{(1)} & \cdot & \cdot & \cdot & x_n^{(1)} \\ x_1^{(2)} & x_2^{(2)} & \cdot & \cdot & \cdot & x_n^{(2)} \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ x_1^{(m)} & x_2^{(m)} & \cdot & \cdot & \cdot & x_n^{(m)} \end{bmatrix} \quad \text{Labels: } \mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \cdot \\ \cdot \\ \cdot \\ y^{(m)} \end{bmatrix}$$

m examples, n independent variables

Forward Pass / Prediction

Id (Model)	Year	Mileage	MPG	Buying Price
5 (Ford Edge)	2018	50,000	23	1 (\$19,000)
105 (Toyota Landcruiser)	2020	10,000	15	0 (\$80,000)
233 (VW Golf)	2010	150,000	36	1 (\$7,000)



Forward Pass / Prediction

Input: \mathbf{X}


Weights: $\mathbf{w} \in \mathbb{R}^{n_x}$, bias: $b \in \mathbb{R}$

$$\mathbf{w} = [w_1 \quad w_2 \quad . \quad . \quad . \quad w_n]$$

Output: $\hat{\mathbf{y}} = \sigma(\mathbf{X}\mathbf{w}^T + b)$

$$\sigma \left(\begin{bmatrix} x_1^{(1)} & x_2^{(1)} & . & . & . & x_n^{(1)} \\ x_1^{(2)} & x_2^{(2)} & . & . & . & x_n^{(2)} \\ . & . & . & . & . & . \\ . & . & . & . & . & . \\ . & . & . & . & . & . \\ x_1^{(m)} & x_2^{(m)} & . & . & . & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} w_1 \\ w_2 \\ . \\ . \\ . \\ w_n \end{bmatrix} + \begin{bmatrix} b \\ b \\ . \\ . \\ . \\ b \end{bmatrix} \right) = \begin{bmatrix} \hat{y}^{(1)} \\ \hat{y}^{(2)} \\ . \\ . \\ . \\ \hat{y}^{(m)} \end{bmatrix}$$

Sigmoid function: $\sigma(z) = \frac{1}{1 + e^{-z}}$



Loss Function


Dataset: $\left\{ \left(\mathbf{x}^{(1)}, y^{(1)} \right), \left(\mathbf{x}^{(2)}, y^{(2)} \right), \dots, \left(\mathbf{x}^{(m)}, y^{(m)} \right) \right\}$

Cross entropy loss function: $\mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) = -(\mathbf{y} \log \hat{\mathbf{y}} + (1 - \mathbf{y}) \log (1 - \hat{\mathbf{y}}))$


Mean squared error (MSE) loss function: $\mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^2$

Cost function: $J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$

Gradient Descent



Find w and b that minimize $J(w, b)$



Gradient Descent

repeat until converge {

$$\mathbf{w} := \mathbf{w} - \alpha \frac{\partial J}{\partial \mathbf{w}}$$

$$b := b - \alpha \frac{\partial J}{\partial b}$$

}

α is the learning rate

